

## Aim: Perform Data Modeling.

### Theory:

**Data Partitioning:** Splitting data into training and testing sets ensures reliable model evaluation. A common split is 75% for training and 25% for testing.

**Visualization:** Bar charts and pie charts help confirm the correct proportions of training and test data.

**Record Count:** Checking the number of records ensures accurate partitioning.

**Two-Sample Z-Test:** This test compares the means of training and test sets to check if they are statistically similar.

**Significance Testing:** If no significant difference is found, the split is unbiased and maintains data integrity.

### Problem Statement:

a. Partition the data set, for example 75% of the records are included in the training data set and 25% are included in the test data set.

## 1.Dataset Partitioning:

- The dataset is split into **75% training data** and **25% test data** using sklearn
- A table or output from a Python script likely shows the total records and their distribution.

```
import pandas as pd

# Load the dataset (Change file name as needed)
file_path = "cleaned_data.csv" # Replace with your actual dataset file
df = pd.read_csv(file_path)

# Display the first few rows to verify the dataset
print(df.head())
```

0	0.000000	0.949	1.000
1	0.000000	0.830	0.961
2	0.529412	0.727	0.952
3	0.529412	0.795	0.938
4	0.529412	0.622	0.867

	Acceptable Streets % - Previous Month \
0	0.985
1	0.983
2	1.000
3	0.850
4	0.800

	Acceptable Sidewalks % - Previous Month \
0	1.000
1	1.000
2	1.000
3	0.941
4	0.973

	Acceptable Streets % - Previous Year \
0	1.00
1	1.00
2	0.84
3	0.96
4	0.88

```
from sklearn.model_selection import train_test_split

# Splitting the dataset into 75% training and 25% testing
train_df, test_df = train_test_split(df, test_size=0.25, random_state=42)

# Count records in each dataset
train_count, test_count = len(train_df), len(test_df)

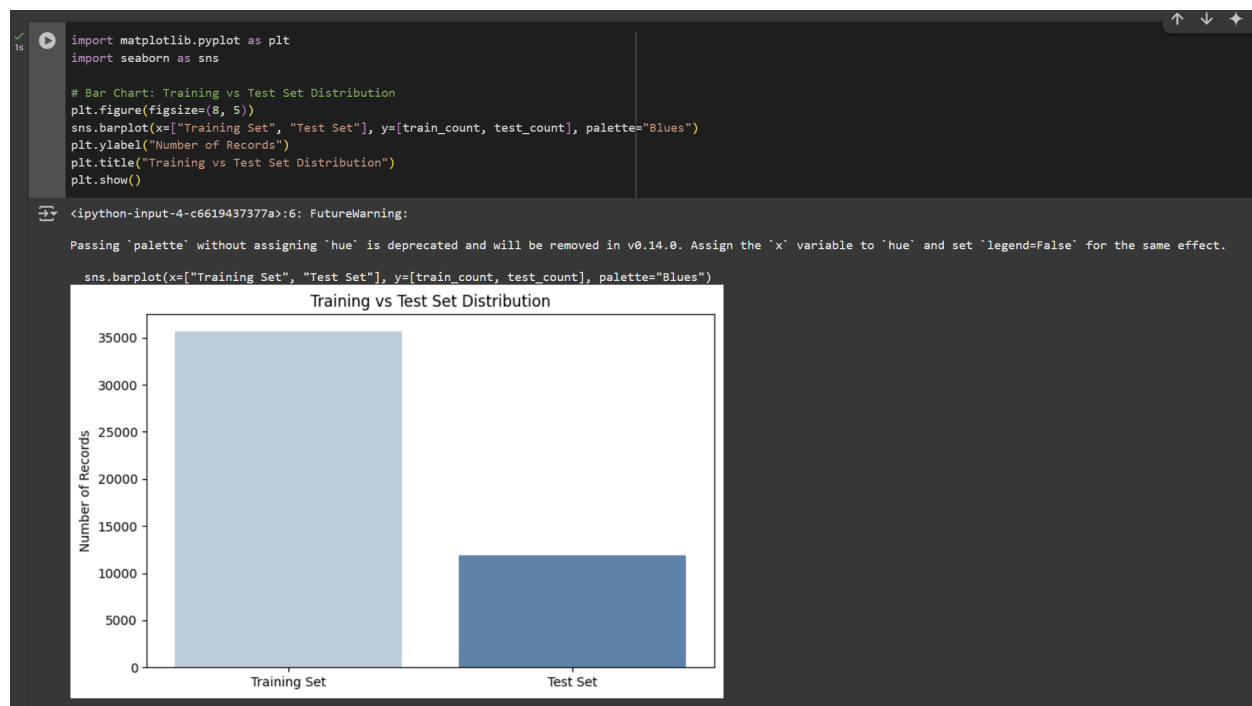
# Print the record count
print(f"Training Set Records: {train_count}")
print(f"Test Set Records: {test_count}")
```

Training Set Records: 35676  
Test Set Records: 11893

b. Use a bar graph and other relevant graph to confirm your proportions.

## 2.Graphs:

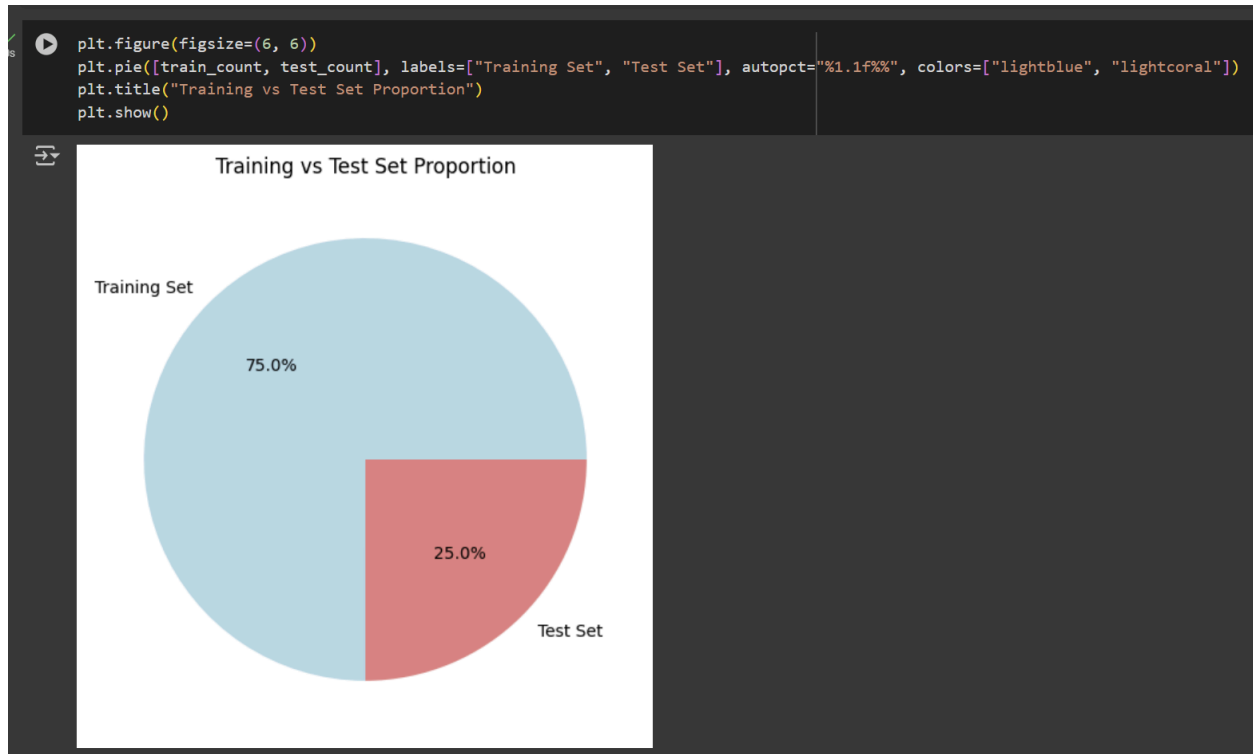
- A **bar graph** and a pie chart is used to verify that the dataset is split correctly.
- The x-axis represents **data categories (training/test)**, while the y-axis represents the **count of records**.



c. Identify the total number of records in the training data set.

### 3.Total Number of Records in the Training Set:

- A numerical output confirms how many records are in the training set.
- This is crucial for ensuring the dataset is correctly partitioned.



d. Validate partition by performing a two-sample Z-test.

### 4.Two-Sample Z-Test for Validation:

- A **statistical test (Z-test)** is performed to validate that the split is unbiased.
- Results likely include **Z-score**, **p-value**, and an interpretation of whether the training and test datasets differ significantly.

```
from scipy import stats

# Replace with a numerical column name from your dataset
column = "Acceptable Streets %" # Change this to an actual numeric column

# Ensure the column exists in the dataset
if column in df.columns:
    print(f"Column '{column}' found. Proceeding with Z-test.")
else:
    print(f"Error: Column '{column}' not found! Choose a valid numeric column.")
```

→ Column 'Acceptable Streets %' found. Proceeding with Z-test.

```
[7] train_values = train_df[column].dropna()
test_values = test_df[column].dropna()

# Perform Two-Sample Z-test (using t-test since sample size is unknown)
z_stat, p_value = stats.ttest_ind(train_values, test_values, equal_var=False)

print(f"Z-statistic: {z_stat}, P-value: {p_value}")

# Interpretation of Z-test results
if p_value > 0.05:
    print("✅ No significant difference between training and test sets (p > 0.05).")
else:
    print("⚠️ Significant difference detected (p ≤ 0.05). Consider re-sampling.")
```

→ Z-statistic: -0.8139349929763461, P-value: 0.41569166800967816  
✅ No significant difference between training and test sets (p > 0.05).

```
import pandas as pd
import numpy as np
from scipy import stats
from sklearn.model_selection import train_test_split

# Load dataset
file_path = "cleaned_data.csv" # Change this to your actual file
df = pd.read_csv(file_path)

# Split the dataset (75% training, 25% test)
train_df, test_df = train_test_split(df, test_size=0.25, random_state=42)

# Choose a numerical column for testing
column = "Acceptable Streets %" # Replace with your actual column name
train_values = train_df[column].dropna() # Remove NaN values

# Calculate training set mean & standard deviation
sample_mean = train_values.mean()
sample_std = train_values.std()
sample_size = len(train_values)

# Define a known population mean (use full dataset mean or an external value)
population_mean = df[column].mean() # You can also use an external reference value

# Calculate Z-score
z_score = (sample_mean - population_mean) / (sample_std / np.sqrt(sample_size))

# Get p-value
p_value = stats.norm.sf(abs(z_score)) * 2 # Two-tailed test

# Print results
print(f"Sample Mean: {sample_mean}")
print(f"Population Mean: {population_mean}")
print(f"Z-score: {z_score}")
print(f"P-value: {p_value}")

# Interpretation
if p_value > 0.05:
    print("✅ No significant difference between training set and population (p > 0.05).")
else:
    print("⚠️ Significant difference detected (p ≤ 0.05).")
```

```
Sample Mean: 0.9319211879134432
Population Mean: 0.9320973112741493
Z-score: -0.405818199815634
P-value: 0.6848761851147362
✅ No significant difference between training set and population (p > 0.05).
```

## Conclusion:

- The experiment **successfully partitions** the dataset using sklearn.
- The **bar graph confirms** the correct proportions of training and test sets.

- A **two-sample Z-test validates** the partitioning statistically.
- The experiment ensures that the dataset is ready for **machine learning model training**.