

Experiment No: 5

Aim: Perform Regression Analysis using Scipy and Sci-kit learn.

Theory

1. Linear Regression:

Linear Regression is a supervised learning algorithm used to predict continuous numerical values.

- It assumes a linear relationship between the independent (input) and dependent (target) variables.
- The model minimizes the sum of squared errors to determine the best-fitting line.
- It is evaluated using performance metrics like Mean Squared Error (MSE) and R-squared (R^2).
- It works best when data follows a linear trend but is sensitive to outliers.

2. Logistic Regression:

Logistic Regression is a classification algorithm used for predicting binary outcomes (0 or 1).

- It applies the sigmoid function to convert linear outputs into probabilities.
- A threshold (typically 0.5) is used to classify data points.
- The model is trained using gradient descent to minimize the log loss function.
- Performance is assessed using accuracy, precision, recall, F1-score, and confusion matrices.

1.Import necessary libraries and load the dataset

```
+ Code + Text

import pandas as pd

# Load the dataset
file_path = "/content/Loan_default.csv"
df = pd.read_csv(file_path)

# Display basic info
print(df.head()) # Show first few rows
print(df.info()) # Show dataset summary
```

2.Check for missing values and anomalies in the dataset.

```
LoanID Age Income LoanAmount CreditScore MonthsEmployed
0 I38PQUQS96 56 85994 50587 520 80
1 HPSK72WA7R 69 50432 124440 458 15
2 C10Z6DPJ8Y 46 84208 129188 451 26
3 V2KKSFM3UN 32 31713 44799 743 0
4 EY08JDHTZP 60 20437 9139 633 8

NumCreditLines InterestRate LoanTerm DTIRatio Education \
0 4 15.23 36 0.44 Bachelor's
1 1 4.81 60 0.68 Master's
2 3 21.17 24 0.31 Master's
3 3 7.07 24 0.23 High School
4 4 6.51 48 0.73 Bachelor's

EmploymentType MaritalStatus HasMortgage HasDependents LoanPurpose
0 Full-time Divorced Yes Yes Other
1 Full-time Married No No Other
2 Unemployed Divorced Yes Yes Auto
3 Full-time Married No No Business
4 Unemployed Divorced No Yes Auto

HasCoSigner Default
0 Yes 0
1 Yes 0
2 No 1
3 No 0
4 No 0
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 255347 entries, 0 to 255346
Data columns (total 18 columns):
```

a) Perform Logistic regression to find out relation between variables

3. Train the dataset using Linear Regression

```
+ Code + Text

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, mean_squared_error

# Load dataset
file_path = "/content/Loan_default.csv"
df = pd.read_csv(file_path)

# Drop LoanID as it's just an identifier
df.drop(columns=["LoanID"], inplace=True)

# Encode categorical variables
categorical_cols = ["Education", "EmploymentType", "MaritalStatus", "HasMortgage",
                    "HasDependents", "LoanPurpose", "HasCosigner"]

label_encoders = {}
for col in categorical_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le

# Split data into features and target
X = df.drop(columns=["Default"])
y = df["Default"]

# Split into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

# (a) Perform Logistic Regression
log_reg = LogisticRegression(max_iter=1000, random_state=42)
log_reg.fit(X_train, y_train)

y_pred = log_reg.predict(X_test)

# Logistic Regression Evaluation
print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```

4. Train the dataset using Logistic Regression and compute:

- Accuracy Score
- Precision, Recall, and F1-score

```
Logistic Regression Accuracy: 0.8854709222635598
Confusion Matrix:
[[45028  111]
 [ 5738  193]]
Classification Report:
              precision    recall  f1-score   support

     0       0.89         1.00         0.94       45139
     1       0.63         0.03         0.06        5931

 accuracy          0.89          0.89          0.89       51070
 macro avg         0.76         0.52         0.50       51070
 weighted avg      0.86         0.89         0.84       51070
```

b) Apply regression model technique to predict the data on the above dataset.

- Find Mean Squared Error (MSE) and R-squared score.

```
reg_model = DecisionTreeRegressor(random_state=42)
reg_model.fit(X_train, y_train)

y_pred_reg = reg_model.predict(X_test)

# Regression Model Evaluation
mse = mean_squared_error(y_test, y_pred_reg)
print("Decision Tree Regression MSE:", mse)
```

Decision Tree Regression MSE: 0.19674955942823574

```
from sklearn.metrics import r2_score

y_prob = log_reg.predict_proba(X_test)[:, 1]
r_squared = r2_score(y_test, y_prob)

print("R-squared value:", r_squared)
```

R-squared value : 0.3017

Conclusion:

The logistic regression model achieves 88.54% accuracy, but its recall for loan defaulters (Class 1) is very low (0.03), indicating poor detection of actual defaults. The decision tree regression model has an MSE of 0.1967 and an R^2 score of 0.3017, showing moderate predictive power. Improving class balance and feature selection could enhance model performance.

Name: Nikita Thadani Class : D15C Roll No: 55